

Exploratory Data Analysis with Pandas

Whenever you receive a new dataset, you should perform some form of EDA. The main things you should be doing are:

- getting a feel for what the data look like
- establishing what types the data have
- identifying any potential issues that need cleaning
- considering whether there may be any issues when it comes to modelling the dataset

For this example, we haven't given you an explicit modelling task to think of whilst you perform this task. This can happen in real life too - "what can you do with this data?" is a legitimate question you may receive. Think about what you might be able to model from this data whilst you get to know the dataset. How might your model be useful to the dataset's owner?

Module load

Please import the libraries you'll need for data analysis: `numpy` , `pandas` , `matplotlib.pyplot` , and `seaborn`

In [1]:

```
# Import the required modules
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

The below is not recommended in general, but will clean up this notebook (there are some issues with seaborn plots kicking out warnings which are a bit messy). Feel free to comment this out if you want to see the warnings.

In [2]:

```
import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)
```

Loading

Load `retail_data.csv` using the relevant pandas method. Make "CustomerID" the index of the DataFrame. Print the first few lines to check your import looks ok.

In [3]:

```
# Import the data
pd.read_csv('data/retail_data.csv')
#notice that the customerID is read in so we could use this as an index, to do this instead
#code, we can use the code below
```

Out[3]:

| | CustomerID | Country | balance | max_spent | mean_spent | min_spent | n_orders | time_between_orders |
|---|------------|----------------|---------|-----------|--------------|-----------|----------|---------------------|
| 0 | 12346 | United Kingdom | 0.00 | 77183.60 | 38591.800000 | 0.00 | 2.0 | |
| 1 | 12348 | Finland | 3874.60 | 2248.80 | 1291.533333 | 478.80 | 3.0 | 54.500000 |
| 2 | 12350 | Norway | 294.40 | 294.40 | 294.400000 | 294.40 | 1.0 | |
| 3 | 12352 | Norway | 1845.13 | 1054.10 | 393.092000 | 0.00 | 5.0 | 11.333333 |
| 4 | 12354 | Spain | 1079.40 | 1079.40 | 1079.400000 | 1079.40 | 1.0 | |
| 5 | 12356 | Portugal | 6621.63 | 5011.34 | 3310.815000 | 1610.29 | 2.0 | 80.000000 |
| 6 | 12358 | Austria | 404.86 | 404.86 | 404.860000 | 404.86 | 1.0 | |
| 7 | 12360 | Austria | 4359.34 | 2984.60 | 2179.670000 | 1374.74 | 2.0 | 88.000000 |

In [4]:

```
pd.read_csv('data/retail_data.csv', index_col='CustomerID')
```

Out[4]:

| | Country | balance | max_spent | mean_spent | min_spent | n_orders | time_between_orders |
|------------|----------------|---------|-----------|--------------|-----------|----------|---------------------|
| CustomerID | | | | | | | |
| 12346 | United Kingdom | 0.00 | 77183.60 | 38591.800000 | 0.00 | 2.0 | NaN |
| 12348 | Finland | 3874.60 | 2248.80 | 1291.533333 | 478.80 | 3.0 | 54.500000 |
| 12350 | Norway | 294.40 | 294.40 | 294.400000 | 294.40 | 1.0 | NaN |
| 12352 | Norway | 1845.13 | 1054.10 | 393.092000 | 0.00 | 5.0 | 11.333333 |
| 12354 | Spain | 1079.40 | 1079.40 | 1079.400000 | 1079.40 | 1.0 | NaN |
| 12356 | Portugal | 6621.63 | 5011.34 | 3310.815000 | 1610.29 | 2.0 | 80.000000 |
| 12358 | Austria | 404.86 | 404.86 | 404.860000 | 404.86 | 1.0 | NaN |

In [5]:

```
#now let us assign the code above to some variables such as
customers = pd.read_csv('data/retail_data.csv', index_col='CustomerID')
```

In [6]:

```
# Print a few rows of the DataFrame
customers.head()
```

Out[6]:

| | Country | balance | max_spent | mean_spent | min_spent | n_orders | time_between_o |
|------------|----------------|---------|-----------|--------------|-----------|----------|----------------|
| CustomerID | | | | | | | |
| 12346 | United Kingdom | 0.00 | 77183.6 | 38591.800000 | 0.0 | 2.0 | |
| 12348 | Finland | 3874.60 | 2248.8 | 1291.533333 | 478.8 | 3.0 | 54.50 |
| 12350 | Norway | 294.40 | 294.4 | 294.400000 | 294.4 | 1.0 | |
| 12352 | Norway | 1845.13 | 1054.1 | 393.092000 | 0.0 | 5.0 | 11.30 |
| 12354 | Spain | 1079.40 | 1079.4 | 1079.400000 | 1079.4 | 1.0 | |

Now what do we have - we have some sort of customer data, we have: the country they are from, a balance, how much they spent, number of orders. So, it looks like this might be some sort of aggregated dataset. I've explicitly not really told you much about this because this is what EDA is about

Understanding the data

The first thing to do is to check the import has definitely worked ok and to get an overview:

1. Get the shape of the data
2. Print some basic statistics about the distribution of each feature e.g. mean, standard deviation, or nr unique values etc. (hint: there is a DataFrame method for this)
3. Get the types of all the data. Do they look right? (answer: they're not strictly right, think about why. However, you don't actually need to edit their types for all the following functions to work)
4. Check missing values in the data - can and should they be cleaned? How would this be handled during the modelling stage?

Now that we are beginning to understand our data we want to do things like: 1.get the shape, 2.basic statistics using the describe method, 3.check data types 4.Are there any rows which contain missing data? Print them (hint: use the methods isnull() and any() find all the rows which contain one or more missing values).

In [7]:

```
# shape
customers.shape
```

#the output shows there are 3254 rows and 11 features

Out[7]:

(3254, 11)

In [8]:

```
# stats
#customers.describe()

#be careful if you use customers.describe().shape
#this will not output all the column, instead we get (8, 10)
#if we use customers.describe(include='all').shape
#we get (11, 11)

#so instead of customers.describe() we should use
customers.describe(include='all')

#there are some columns that are categorical, for example the column is a categorical vari
```

Out[8]:

| | Country | balance | max_spent | mean_spent | min_spent | n_orders | time_b |
|---------------|----------------|---------------|--------------|--------------|-------------|-------------|--------|
| count | 3254 | 3254.000000 | 3254.000000 | 3254.000000 | 3254.000000 | 3254.000000 | |
| unique | 11 | NaN | NaN | NaN | NaN | NaN | |
| top | United Kingdom | NaN | NaN | NaN | NaN | NaN | |
| freq | 3013 | NaN | NaN | NaN | NaN | NaN | |
| mean | NaN | 3595.679152 | 1241.132434 | 697.714203 | 336.855298 | 4.025814 | |
| std | NaN | 12287.159793 | 2765.253795 | 1105.379255 | 592.529725 | 6.330090 | |
| min | NaN | -1192.200000 | 0.000000 | 0.000000 | 0.000000 | 1.000000 | |
| 25% | NaN | 330.807500 | 298.307500 | 215.720000 | 0.000000 | 1.000000 | |
| 50% | NaN | 948.595000 | 658.430000 | 463.260000 | 150.000000 | 2.000000 | |
| 75% | NaN | 3198.430000 | 1358.170000 | 852.338786 | 418.197500 | 4.000000 | |
| max | NaN | 394689.180000 | 77183.600000 | 38591.800000 | 9885.320000 | 134.000000 | |

In [9]:

```
# data types as imported
customers.dtypes
```

Out[9]:

```
Country          object
balance          float64
max_spent        float64
mean_spent       float64
min_spent        float64
n_orders         float64
time_between_orders float64
total_items      float64
total_items_returned float64
total_refunded   float64
total_spent      float64
dtype: object
```

Are there any rows which contain missing data? Print them (hint: use the methods `isnull()` and `any()` find all the rows which contain one or more missing values).

In [10]:

```
# Print missing value rows
#using:
customers.isnull()
#or use below to check across column
customers.isnull().any(axis=1)

#you can check this by returning the number of rows using
#customers.isnull().any(axis=1).shape    #this returns 3254
```

Out[10]:

```
CustomerID
12346      True
12348     False
12350      True
12352     False
12354      True
12356     False
12358      True
12360     False
12361      True
12362     False
12364      True
12373      True
12377     False
12378      True
12379      True
12380      True
12381      True
12383     False
12384      True
12395     False
12397      True
12399     False
12401      True
12402      True
12405      True
12407     False
12408     False
12409      True
12410     False
12413     False
...
18233     True
18235     False
18236     False
18237     False
18239     False
18241     False
18242     False
18245     False
18246     True
18248     False
18250     True
18252     True
18256     True
18257     False
```

```

18259    True
18260    False
18262     True
18263    False
18265     True
18268     True
18269     True
18270     True
18272    False
18273     True
18277     True
18280     True
18281     True
18282     True
18283    False
18287     True
Length: 3254, dtype: bool

```

There are actually many...but you'll see that this makes sense given the data. Are there any rows that have missing values apart from this obvious column? Rewrite the cell above and write code to check (tip: you can use `.drop`).

In [11]:

```

#we can further index the dataframe using
customers[customers.isnull().any(axis=1)]

```

Out[11]:

| | Country | balance | max_spent | mean_spent | min_spent | n_orders | time_between_orders | t |
|------------|----------------|---------|-----------|--------------|-----------|----------|---------------------|---|
| CustomerID | | | | | | | | |
| 12346 | United Kingdom | 0.00 | 77183.60 | 38591.800000 | 0.00 | 2.0 | NaN | |
| 12350 | Norway | 294.40 | 294.40 | 294.400000 | 294.40 | 1.0 | NaN | |
| 12354 | Spain | 1079.40 | 1079.40 | 1079.400000 | 1079.40 | 1.0 | NaN | |
| 12358 | Austria | 404.86 | 404.86 | 404.860000 | 404.86 | 1.0 | NaN | |
| 12361 | Belgium | 174.90 | 174.90 | 174.900000 | 174.90 | 1.0 | NaN | |
| 12364 | Belgium | 1840.52 | 1840.52 | 1840.520000 | 1840.52 | 1.0 | NaN | |
| 12373 | Austria | 324.60 | 324.60 | 324.600000 | 324.60 | 1.0 | NaN | |

In [12]:

```
# Other missing value columns?  
#yes, if we look at the column time_between_orders that is a significant one, either becau  
#or there is only one order, so we can drop that column  
customers.drop('time_between_orders', axis=1).isnull().any(axis=1)
```

Out[12]:

```
CustomerID  
12346      False  
12348      False  
12350      False  
12352      False  
12354      False  
12356      False  
12358      False  
12360      False  
12361      False  
12362      False  
12364      False  
12373      False  
12377      False  
12378      False  
12379      False  
12380      False  
12381      False  
12383      False  
12384      False  
12395      False  
12397      False  
12399      False  
12401      False  
12402      False  
12405      False  
12407      False  
12408      False  
12409      False  
12410      False  
12413      False  
...  
18233      False  
18235      False  
18236      False  
18237      False  
18239      False  
18241      False  
18242      False  
18245      False  
18246      False  
18248      False  
18250      False  
18252      False  
18256      False  
18257      False  
18259      False  
18260      False  
18262      False  
18263      False
```



```
18265    False
18268    False
18269    False
18270    False
18272    False
18273    False
18277    False
18280    False
18281    False
18282    False
18283    False
18287    False
Length: 3254, dtype: bool
```

In [13]:

```
#again we can index the customer dataset
customers[customers.drop('time_between_orders', axis=1).isnull().any(axis=1)]
```

Out[13]:

| | Country | balance | max_spent | mean_spent | min_spent | n_orders | time_between_or |
|------------|---------|---------|-----------|------------|-----------|----------|-----------------|
| CustomerID | | | | | | | |

As a bonus, write some code which checks if a column looks more like an integer than a float value. Skip this if you can't do it quickly (it's not necessary for what comes next).

In [14]:

```
# BONUS: skip if stuck. Find integer columns
for col in customers.columns:
    col_series = customers[col]
    dtype = col_series.dtype
    has_missing = col_series.isnull().any()
    if dtype == np.dtype('float64') and not has_missing:
        if (col_series.astype('int') == col_series).sum() == col_series.shape[0]:
            print('{} looks like an integer'.format(col))
            print('Converting it to an integer')
            customers[col] = customers[col].astype('int')
customers.dtypes
```

```
n_orders looks like an integer
Converting it to an integer
total_items looks like an integer
Converting it to an integer
total_items_returned looks like an integer
Converting it to an integer
```

Out[14]:

```
Country          object
balance          float64
max_spent         float64
mean_spent        float64
min_spent         float64
n_orders          int32
time_between_orders float64
total_items        int32
total_items_returned int32
total_refunded     float64
total_spent        float64
dtype: object
```

===== QUESTIONS

To be good at EDA, you need to be able to quickly answer simple questions about the data. As you become more proficient at pandas, and using `.groupby`, `.apply`, and other methods, you'll become very fast at answering quick questions. [This is a great link \(https://pandas.pydata.org/pandas-docs/stable/groupby.html\)](https://pandas.pydata.org/pandas-docs/stable/groupby.html) to the pandas docs which outlines methods you'll find useful.

Below are a series of questions which are reasonable first questions to ask about the data. They start simple and increase in difficulty. There are many ways to get the answers and, for that matter, display them. In my sample solutions, I normally opt to plot solutions when it's appropriate and easy to do so. Feel free to print tables, or even write for loops if you prefer!

If you get inspired, ask and answer your own questions.

This is your chance to investigate and get a feel for this dataset.

1) How many unique customers are there? And, therefore, what is the average number of rows per customer?

In [15]:

```
# Nr unique customers
customers.index.nunique()
```

Out[15]:

3254

In [16]:

```
# Rows per customer - the number of rows in the dataset /
customers.index.nunique()/customers.shape[0]
```

Out[16]:

1.0

2) What is the total amount spent by customers? (feel free to display numbers in a nice format)

In [17]:

```
# sum of total spent by all the customers
customers.total_spent.sum() #this returns 11894018.760000002
#to print it out nicely in a scientific format
'{:.2e}'.format(customers.total_spent.sum())
```

Out[17]:

'1.19e+07'

3) What is the total amount refunded to customers?

In [18]:

```
# Total amount refunded can be returned using customers.total_refunded.sum(), however thi
#the refunds are show as negative figures in the dataset. To make it output a positive numi
-customer.total_refunded.sum()
```

Out[18]:

193678.8

4) Assuming the company serving these customers had a balance of 0 at the start, what is their balance now?

In [19]:

```
# The balance of the company right now  
#we can subtract the total refunded from the total spent by customers, but in this case we  
#discussed previously  
customers.total_spent.sum() + customers.total_refunded.sum()
```

Out[19]:

11700339.96


5) Count the number of customers from each country, and all countries excluding the UK

In [20]:

```
# Count of customers by country - use seaborn!  
#to do this we can use a countplot, the output on the x axis is messy so rotate the labels  
sns.countplot(customers.Country)  
plt.xticks(rotation=90)
```

Out[20]:

(array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10]),
<a list of 11 Text xticklabel objects>)




In [21]:

```
# ...excluding UK  
sns.countplot(customers.Country[customers.Country != 'United Kingdom'])  
plt.xticks(rotation=90)
```

Out[21]:

(array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9]), <a list of 10 Text xticklabel objects>)



6) What is the distribution of the number of orders customers have made? If the distribution is hard to visualise or unclear, try splitting the range up somehow

In [22]:


```
# Distribution of number of orders each customer makes - we will use distplot  
sns.distplot(customers['n_orders'])
```

C:\ProgramData\Anaconda3\lib\site-packages\matplotlib\axes_axes.py:6462: UserWarning: The 'normed' kwarg is deprecated, and has been replaced by the 'density' kwarg.

warnings.warn("The 'normed' kwarg is deprecated, and has been "

Out[22]:

<matplotlib.axes._subplots.AxesSubplot at 0x1b560925cc0>



It is not very clear what is going on in that plot. It appears there are some rows with very large values. you might want to have a quick look at all the distribution of all the rows that have some smaller numbers

In [23]:

```
# Distribution of the majority of orders
#we will choose a random number 30 and print a subset of the dataset with the rows with order_id <= 30
#as this is small we can use a countplot to examine the data
x = 'n_orders'
query = 'n_orders <= 30'
df = customers.query(query)
df
```

Out[23]:


| | Country | balance | max_spent | mean_spent | min_spent | n_orders | time_between_orders | 1 |
|------------|----------------|---------|-----------|--------------|-----------|----------|---------------------|---|
| CustomerID | | | | | | | | |
| 12346 | United Kingdom | 0.00 | 77183.60 | 38591.800000 | 0.00 | 2 | NaN | |
| 12348 | Finland | 3874.60 | 2248.80 | 1291.533333 | 478.80 | 3 | 54.500000 | |
| 12350 | Norway | 294.40 | 294.40 | 294.400000 | 294.40 | 1 | NaN | |
| 12352 | Norway | 1845.13 | 1054.10 | 393.092000 | 0.00 | 5 | 11.333333 | |
| 12354 | Spain | 1079.40 | 1079.40 | 1079.400000 | 1079.40 | 1 | NaN | |
| 12356 | Portugal | 6621.63 | 5011.34 | 3310.815000 | 1610.29 | 2 | 80.000000 | |
| 12358 | Austria | 404.86 | 404.86 | 404.860000 | 404.86 | 1 | NaN | |

In [24]:

```
#using countplot to plot the table above  
sns.countplot(df[x])  
plt.xticks(rotation=90)
```

Out[24]:

```
(array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15, 16,  
        17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29]),  
<a list of 30 Text xticklabel objects>)
```



In [25]:

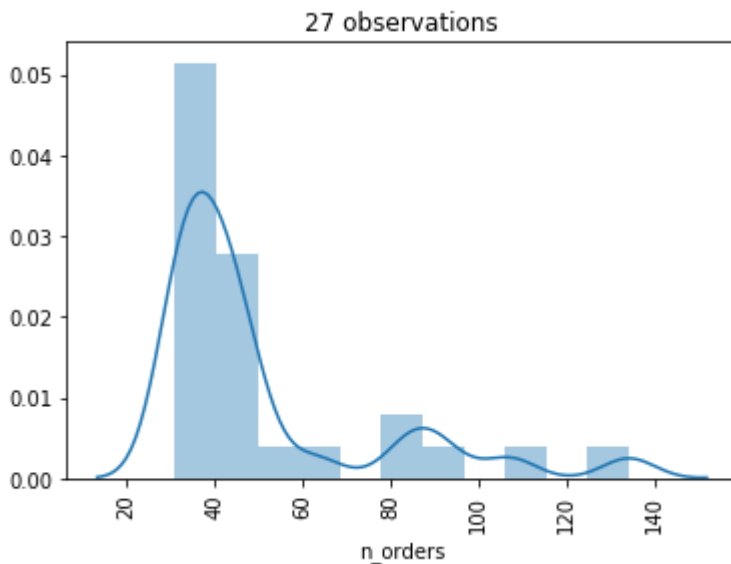
```
# Distribution of the outliers
#Here we could look at orders that are greater than 30 to get the outliers. Instead of a c
#as the data is more disparate. from the plot we see that there are 27 observations that h
x = 'n_orders'
query = 'n_orders > 30'
df = customers.query(query)
sns.distplot(df[x])
plt.xticks(rotation=90)
plt.title('{} observations'.format(df.shape[0]))
```

C:\ProgramData\Anaconda3\lib\site-packages\matplotlib\axes_axes.py:6462: UserWarning: The 'normed' kwarg is deprecated, and has been replaced by the 'density' kwarg.

warnings.warn("The 'normed' kwarg is deprecated, and has been "

Out[25]:

Text(0.5,1,'27 observations')



7) What is the distribution of the amount spent by customers? Again, if the distribution is hard to visualise, try splitting the range up somehow. Be careful to show that you understand how many observations you are dealing with in each plot.

In [26]:


```
# Distribution of the total spent  
sns.distplot(customers['total_spent'])
```

C:\ProgramData\Anaconda3\lib\site-packages\matplotlib\axes_axes.py:6462: UserWarning: The 'normed' kwarg is deprecated, and has been replaced by the 'density' kwarg.

warnings.warn("The 'normed' kwarg is deprecated, and has been "

Out[26]:

<matplotlib.axes._subplots.AxesSubplot at 0x1b561efa908>



In [27]:

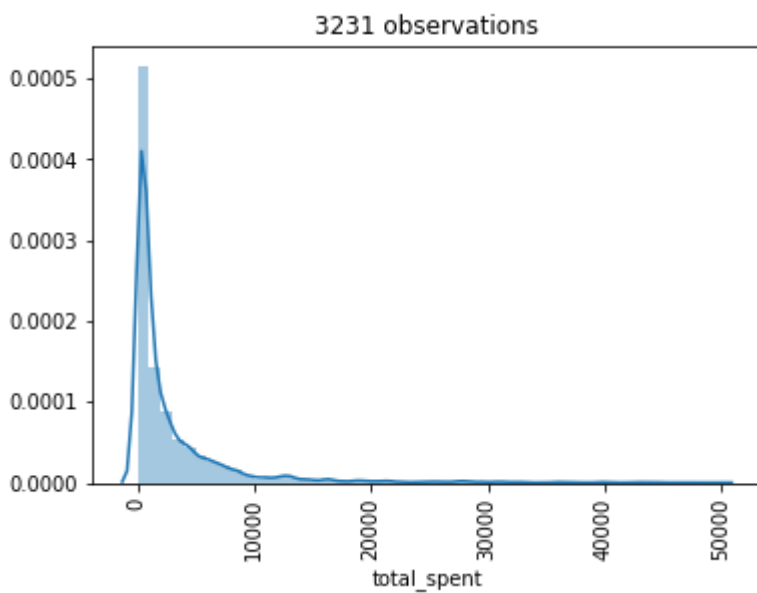
```
# ...the distribution for the majority of companies
x = 'total_spent'
query = 'total_spent <= 50000'
df = customers.query(query)
sns.distplot(df[x])
plt.xticks(rotation=90)
plt.title('{} observations'.format(df.shape[0]))
```

C:\ProgramData\Anaconda3\lib\site-packages\matplotlib\axes_axes.py:6462: UserWarning: The 'normed' kwarg is deprecated, and has been replaced by the 'density' kwarg.

warnings.warn("The 'normed' kwarg is deprecated, and has been "

Out[27]:

Text(0.5,1,'3231 observations')



In [28]:

```
#the above plot shows that there are 3231 observations with total spent over 50,000. how many
#dataset. There are 3254 so the lion share is over 500000
customers.shape[0]
```

Out[28]:

3254

In [29]:

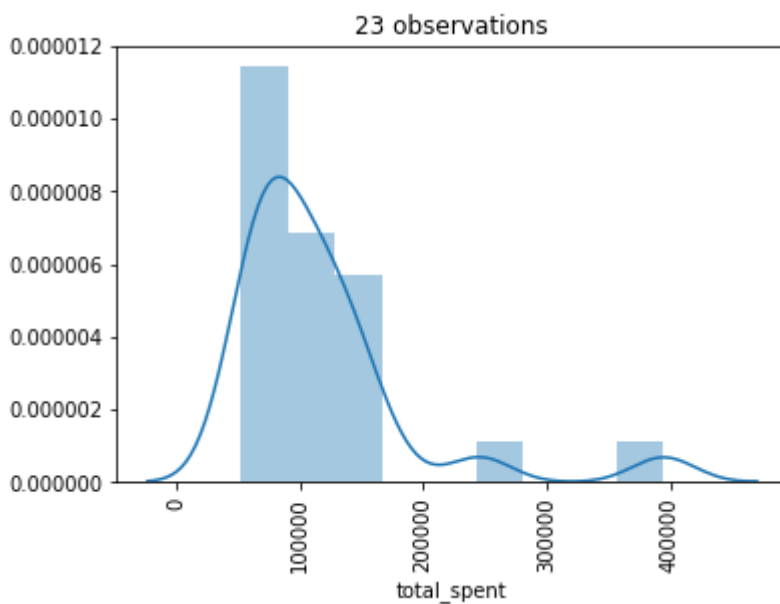
```
# ...the distribution for the outliers
x = 'total_spent'
query = 'total_spent > 50000'
df = customers.query(query)
sns.distplot(df[x])
plt.xticks(rotation=90)
plt.title('{} observations'.format(df.shape[0]))
```

C:\ProgramData\Anaconda3\lib\site-packages\matplotlib\axes_axes.py:6462: UserWarning: The 'normed' kwarg is deprecated, and has been replaced by the 'density' kwarg.

warnings.warn("The 'normed' kwarg is deprecated, and has been "

Out[29]:

Text(0.5,1,'23 observations')



8) How about the distribution of refunds? Again, if prevalent values or outliers are making it difficult to visualise the distribution, split the range somehow.

In [40]:


```
# Distribution of refunds  
sns.distplot(customers['total_refunded'])
```

C:\ProgramData\Anaconda3\lib\site-packages\matplotlib\axes_axes.py:6462: UserWarning: The 'normed' kwarg is deprecated, and has been replaced by the 'density' kwarg.

warnings.warn("The 'normed' kwarg is deprecated, and has been "

Out[40]:

<matplotlib.axes._subplots.AxesSubplot at 0x1b562113f98>



In [41]:

```
# ...the distribution of the majority
prop_refunded = sum(customers.total_refunded == 0) / customers.shape[0]
print('{:.2f} of customers have never had a refund'.format(prop_refunded))
df = customers.query('total_refunded < 0')
print('Range of the orders with a refund: ({} , {})' .format(
    df['total_refunded'].min(), df['total_refunded'].max()))
query = 'total_refunded < 0 and total_refunded > -1000'
df = customers.query(query)
sns.distplot(df['total_refunded'], rug=False)
plt.title('{} observations {}'.format(df.shape[0], query))
plt.show()
```


0.66 of customers have never had a refund

Range of the orders with a refund: (-77183.6, -0.42)

C:\ProgramData\Anaconda3\lib\site-packages\matplotlib\axes_axes.py:6462: UserWarning: The 'normed' kwarg is deprecated, and has been replaced by the 'density' kwarg.

warnings.warn("The 'normed' kwarg is deprecated, and has been "

1091 observations total_refunded < 0 and total_refunded > -1000



In [42]:


```
# ...the distribution for the outliers
query = 'total_refunded <= -1000'
df = customers.query(query)
sns.distplot(df['total_refunded'], rug=False)
plt.title('{} observations {}'.format(df.shape[0], query))
```

C:\ProgramData\Anaconda3\lib\site-packages\matplotlib\axes_axes.py:6462: UserWarning: The 'normed' kwarg is deprecated, and has been replaced by the 'density' kwarg.

warnings.warn("The 'normed' kwarg is deprecated, and has been "

Out[42]:

Text(0.5,1,'16 observations total_refunded <= -1000')




9) Use a [violinplot](https://seaborn.pydata.org/generated/seaborn.violinplot.html?highlight=violin#seaborn.violinplot) (<https://seaborn.pydata.org/generated/seaborn.violinplot.html?highlight=violin#seaborn.violinplot>) and/or a [boxplot](https://seaborn.pydata.org/generated/seaborn.boxplot.html) (<https://seaborn.pydata.org/generated/seaborn.boxplot.html>) to plot the distribution of the total spent per country

In [45]:

```
# Violin and box plots of the total spend broken down by country video 8
sns.violinplot(x='Country', y='total_spent', data=customers)
plt.xticks(rotation=90)
plt.ylim(0, 20000)
```

Out[45]:

(0, 20000)




In [47]:

```
sns.boxplot(x='Country', y='total_spent', data=customers)
plt.xticks(rotation=90)
plt.ylim(0, 20000)
```

Out[47]:

(0, 20000)



10a) What is the total amount spent broken down by country? One option here is to use pandas groupby to create a DataFrame containing the required information, then use sns.barplot to plot the information.

In [49]:

```
# Sum of total spent broken down by country
(customers['total_spent']
 .groupby(customers['Country'])
 .sum()
)
#to change the way this is displayed and to add a new column called sum of total spent add
```

Out[49]:

```
Country
Austria          10619.20
Belgium          46252.29
Finland          31091.07
France           272068.47
Germany          325883.91
Italy            20046.71
Norway           34714.18
Portugal         33414.90
Spain            95774.83
Switzerland       67361.13
United Kingdom   10956792.07
Name: total_spent, dtype: float64
```

In [52]:

```
# Sum of total spent broken down by country
(customers['total_spent']
 .groupby(customers['Country'])
 .sum()
 .rename('sum of total spent')
 .reset_index()
)
```

Out[52]:

| | Country | sum of total spent |
|----|----------------|--------------------|
| 0 | Austria | 10619.20 |
| 1 | Belgium | 46252.29 |
| 2 | Finland | 31091.07 |
| 3 | France | 272068.47 |
| 4 | Germany | 325883.91 |
| 5 | Italy | 20046.71 |
| 6 | Norway | 34714.18 |
| 7 | Portugal | 33414.90 |
| 8 | Spain | 95774.83 |
| 9 | Switzerland | 67361.13 |
| 10 | United Kingdom | 10956792.07 |

In [57]:

```
#to do the plot below we will assign the code above to a variable
# Sum of total spent broken down by country
sum_tot_spent_by_country = (customers['total_spent']
    .groupby(customers['Country'])
    .sum()
    .rename('sum of total spent')
    .reset_index()


)
```

In [58]:

```
# Plot the data
sns.barplot(x='Country', y='sum of total spent', data=sum_tot_spent_by_country)
plt.xticks(rotation=90)
```

Out[58]:

```
(array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10]),
 <a list of 11 Text xticklabel objects>)
```



10b) Again, as might be expected, there is one dominant country. Plot the data again, but exclude the dominant country. In the title of the graph, write explicitly what proportion of the total spent is shown in the plot.

In [61]:

```
# Plot of total spend broken down by country, excluding UK
df = sum_tot_spent_by_country.query('Country != "United Kingdom"')
df
#to calculate the sum and plot this information, see below code
```

Out[61]:


| | Country | sum of total spent |
|---|-------------|--------------------|
| 0 | Austria | 10619.20 |
| 1 | Belgium | 46252.29 |
| 2 | Finland | 31091.07 |
| 3 | France | 272068.47 |
| 4 | Germany | 325883.91 |
| 5 | Italy | 20046.71 |
| 6 | Norway | 34714.18 |
| 7 | Portugal | 33414.90 |
| 8 | Spain | 95774.83 |
| 9 | Switzerland | 67361.13 |

In [62]:

```
# Plot of total spend broken down by country, excluding UK
df = sum_tot_spent_by_country.query('Country != "United Kingdom"')
sns.barplot(x='Country', y='sum of total spent', data=df)
plt.xticks(rotation=90)
prop_uk_spending = customers.query('Country == "United Kingdom"').total_spent.sum() / cust
plt.title('{:.3f}'.format(1-prop_uk_spending))
```

Out[62]:

Text(0.5,1,'0.079')



10c) What is the average amount spent per country? (Hint: you'll be able to use very similar code to the analysis summing the total spent)


In [63]:

```
# Mean total_spent by country
mean_tot_spent_by_country = (customers['total_spent']
    .groupby(customers['Country'])
    .mean()
    .rename('mean of total spent')
    .reset_index()

)
sns.barplot(x='Country', y='mean of total spent', data=mean_tot_spent_by_country)
plt.xticks(rotation=90)
```

Out[63]:

```
(array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10]),
 <a list of 11 Text xticklabel objects>)
```



10d) Bonus: (move on if not solved in a few mins) -- What is the average amount spent *per order* per country (there is a column `n_orders`)? Hint: be careful...the mean of the column `mean_spent` per

country is not the same as the desired quantity e.g. $\frac{\frac{10}{5} + \frac{20}{4}}{2} \neq \frac{10+20}{5+4}$...my solution first makes a function which calculates the mean spend per order for a given dataframe, then I use `apply` to apply this to each country using a `groupby`.

In [64]:

```
# Mean amount spent per order broken down by country -
#You will not get the correct averages if you sum up the totals from the column n_orders and
#you need to do something more complicated which will be shown in the cell below
customers
```

Out[64]:

| | Country | balance | max_spent | mean_spent | min_spent | n_orders | time_between_orders | 1 |
|------------|----------------|---------|-----------|--------------|-----------|----------|---------------------|---|
| CustomerID | | | | | | | | |
| 12346 | United Kingdom | 0.00 | 77183.60 | 38591.800000 | 0.00 | 2 | NaN | |
| 12348 | Finland | 3874.60 | 2248.80 | 1291.533333 | 478.80 | 3 | 54.500000 | |
| 12350 | Norway | 294.40 | 294.40 | 294.400000 | 294.40 | 1 | NaN | |
| 12352 | Norway | 1845.13 | 1054.10 | 393.092000 | 0.00 | 5 | 11.333333 | |
| 12354 | Spain | 1079.40 | 1079.40 | 1079.400000 | 1079.40 | 1 | NaN | |
| 12356 | Portugal | 6621.63 | 5011.34 | 3310.815000 | 1610.29 | 2 | 80.000000 | |
| 12358 | Austria | 404.86 | 404.86 | 404.860000 | 404.86 | 1 | NaN | |

In [69]:

```

#we will define a function to do the above
def mean_spent_per_order(df):
    total_orders = df.n_orders.sum()
    total_spent = df.total_spent.sum()
    return total_spent / total_orders    #returns the average per order
mean_spent_per_order(customers)
#breakdown by country
avg_order_spent_by_country = (
    customers
    .groupby('Country')
    .apply(mean_spent_per_order)
    .rename('mean spend per order')
    .reset_index()
)
sns.barplot(x='Country', y='mean spend per order', data=avg_order_spent_by_country)
plt.xticks(rotation=90)


```

Out[69]:

```

(array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10]),
 <a list of 11 Text xticklabel objects>)

```



11a) Since the total number of orders is very skewed, when comparing countries, we should compare proportions. Compare the proportion of customers with 1 order vs. 2+ orders for each country. Put this information in one table, or one plot if you can.

In [74]:

```

df = customers
prop_df = (df['n_orders']
            .apply(lambda x: '1' if x==1 else '>2')
            .rename('Number of Orders')
            .groupby(df['Country'])
            .value_counts(normalize=True)
            .rename('prop')
            .reset_index()
)
sns.barplot(x='Country', y='prop', hue='Number of Orders', data=prop_df)
plt.xticks(rotation=90)


```

Out[74]:

```

(array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10]),
<a list of 11 Text xticklabel objects>)

```




11b) Use `pd.cut` to repeat the above but with some more interesting range bins

In [77]:

```
#we wil copy the function from above and replace lines of it with pd.cut, to see how pd.cut
#cell below and the output
df = customers
prop_df = (pd.cut(df['n_orders'], [0, 1, 2, 5, 10, np.inf])
            .rename('Number of Orders')
            .groupby(df['Country'])
            .value_counts(normalize=True)
            .rename('prop')
            .reset_index()
            )
sns.barplot(x='Country', y='prop', hue='Number of Orders', data=prop_df)
plt.xticks(rotation=90)
```

Out[77]:

```
(array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10]),
 <a list of 11 Text xticklabel objects>)
```



In [75]:

```
pd.cut(df['n_orders'], [0, 1, 2, 5, 10, np.inf])
```

Out[75]:

```
CustomerID
12346      (1.0, 2.0]
12348      (2.0, 5.0]
12350      (0.0, 1.0]
12352      (2.0, 5.0]
12354      (0.0, 1.0]
12356      (1.0, 2.0]
12358      (0.0, 1.0]
12360      (1.0, 2.0]
12361      (0.0, 1.0]
12362      (5.0, 10.0]
12364      (0.0, 1.0]
12373      (0.0, 1.0]
12377      (1.0, 2.0]
12378      (0.0, 1.0]
12379      (1.0, 2.0]
12380      (0.0, 1.0]
12381      (0.0, 1.0]
12383      (5.0, 10.0]
12384      (0.0, 1.0]
12395      (5.0, 10.0]
12397      (0.0, 1.0]
12399      (2.0, 5.0]
12401      (0.0, 1.0]
12402      (0.0, 1.0]
12405      (0.0, 1.0]
12407      (2.0, 5.0]
12408      (5.0, 10.0]
12409      (1.0, 2.0]
12410      (2.0, 5.0]
12413      (2.0, 5.0]
...
18233      (0.0, 1.0]
18235      (1.0, 2.0]
18236      (2.0, 5.0]
18237      (1.0, 2.0]
18239      (2.0, 5.0]
18241      (10.0, inf]
18242      (1.0, 2.0]
18245      (5.0, 10.0]
18246      (0.0, 1.0]
18248      (2.0, 5.0]
18250      (1.0, 2.0]
18252      (0.0, 1.0]
18256      (0.0, 1.0]
18257      (5.0, 10.0]
18259      (0.0, 1.0]
18260      (5.0, 10.0]
18262      (0.0, 1.0]
18263      (1.0, 2.0]
18265      (0.0, 1.0]
18268      (1.0, 2.0]
18269      (1.0, 2.0]
18270      (1.0, 2.0]
```

```

18272    (2.0, 5.0]
18273    (0.0, 1.0]
18277    (0.0, 1.0]
18280    (0.0, 1.0]
18281    (0.0, 1.0]
18282    (1.0, 2.0]
18283    (5.0, 10.0]
18287    (0.0, 1.0]

```

Name: n_orders, Length: 3254, dtype: category

Categories (5, interval[float64]): [(0.0, 1.0] < (1.0, 2.0] < (2.0, 5.0] < (5.0, 10.0] < (10.0, inf]]

12) In your own words, summarise what you've found out about this dataset from your analysis above

In [39]:

```

# Interpret your findings above!
# Interpret your findings above!
# * This is a dataset about orders customers have made with a business
# * The data are aggregated up to the customer level
# * The features are summary statistics about their order history
# -----
# * There is a total spending of 11 million in this dataset (information about
#   what denomination is not contained in the dataset - it could be cents!)
# * There is comparatively little refunded
# * The customers are mostly from the UK
# * Of customers excluding the UK, France and Germany are most common
# * From simple distribution plots of spending and n_order variables, we can
#   see we have some companies that are big outliers
# * The pattern is roughly the same regarding the total amount spent
# * The distribution of total spent by customers from each region, the
#   distributions aren't radically different...but there are a few outliers
#   with very large spends
# * The mean order is about £1000 (and the mean item cost is a mere £1.69...
#   my bet is that this dataset is about selling office stationary at this
#   point)
# * For each country, between 30 and 50 percent of customers have submitted
#   only one order (Bonus: germany seems to have loyal customers, over 10%
#   of the customers have ordered 10 times or more!)
# * ...this is by no means an exhaustive list!!!

```